# Big Data (ECE 595-004/007), Fall 2018
## Hands-on 1, Part2: Hadoop Set-up (15 pts)
### Due date: September 21, 2018, 11:59 PM

This assignment is straightforward. You have to set-up Hadoop in your virtual machine. For grading, submit the screenshot of the steps that you followed for set-up in one word file and upload the PDF for it.

**Virtual machine**

The tutorial for VMWare player has been given in the tutorial document.
Those who have Mac, they can install Virtual Box and create a virtual machine using it.
Link for Virtual Box installation: https://www.youtube.com/watch?v=lEvM-No4eQo
Download Ubuntu 16: http://releases.ubuntu.com/16.04/ubuntu-16.04.5-desktop-amd64.iso if your machine is 64-bit or http://releases.ubuntu.com/16.04/ubuntu-16.04.5-desktop-i386.iso if your machine is 32-bit.
Create Ubuntu virtual machine: https://www.youtube.com/watch?v=fh8OdDd0K30

**Hadoop set-up**

1. Install Java: **`sudo apt-get install default-jdk`**
   You can confirm the installation**: `java -version`**
2. Install SSH server: **`sudo apt-get install openssh-server`**
3. Download Hadoop 2.9.1 from the following link:
   https://www.apache.org/dyn/closer.cgi/hadoop/common/hadoop-2.9.1/hadoop-2.9.1.tar.gz
4. Go to the Downloads directory: **`cd Downloads`**
5. Extract Hadoop files from the downloaded file: **`tar xzvf hadoop-2.9.1.tar.gz`**
6. Move extracted Hadoop directory to /usr/local directory:
   **`sudo mv hadoop-2.9.1 /usr/local/hadoop`**
   Change ownership of Hadoop directory:
   **`sudo chown –R bigdata:bigdata /usr/local/hadoop`**
   Note: **`bigdata`** is a user here. It may be different in your case.
7. Find directory for Java: **`readlink -f /usr/bin/java | sed "s:/bin/java::"`**
   The command output will be the directory for Java.
8. Go to home directory: **`cd`**
9. Open .bashrc file: **`gedit .bashrc`**
   Add following lines:
   ```
   export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/jre
   export HADOOP_HOME=/usr/local/Hadoop
   export PATH=$PATH:$HADOOP_HOME/bin:$HADOOP_HOME/sbin
   ```
   Note: Value for JAVA_HOME is the output of the previous command for finding Java directory
   Close .bashrc file and execute command: **`. .bashrc`**
   You can verify the changes with the command: **`echo $JAVA_HOME`**
10. Open /usr/local/hadoop/etc/hadoop/hadoop-env.sh:
    **`gedit /usr/local/hadoop/etc/hadoop/hadoop-env.sh`**

Replace **export JAVA_HOME=${JAVA_HOME}** with
**export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64/jre**
Also add this: **export HADOOP_OPTS=-Djava.net.preferIPv4Stack=true**
Close the file

11. Create a directory /app/hadoop/tmp: **sudo mkdir -p /app/hadoop/tmp**
    Change ownership: **sudo chown –R bigdata:bigdata /app/hadoop/tmp**
12. Open /usr/local/hadoop/etc/hadoop/core-site.xml:
    **gedit /usr/local/hadoop/etc/hadoop/core-site.xml**
    Add these lines inside <configuration> </configuration> tags
    **<property>**
    **<name>hadoop.tmp.dir</name>**
    **<value>/app/hadoop/tmp</value>**
    **</property>**
    **<property>**
    **<name>fs.default.name</name>**
    **<value>hdfs://localhost:9000</value>**
    **</property>**
    Close the file
13. Open /usr/local/hadoop/etc/hadoop/hdfs-site.xml:
    **gedit /usr/local/hadoop/etc/hadoop/hdfs-site.xml**
    Add these lines inside <configuration> </configuration> tags
    **<property>**
    **<name>dfs.replication</name>**
    **<value>1</value>**
    **</property>**
    Close the file.
14. Open /usr/local/hadoop/etc/hadoop/mapred-site.xml:
    **gedit /usr/local/hadoop/etc/hadoop/mapred-site.xml**
    Add these lines inside <configuration> </configuration> tags
    **<property>**
    **<name>mapreduce.framework.name</name>**
    **<value>yarn</value>**
    **</property>**
    Note: mapred-site.xml is not available. Rename mapred-site.xml.template file to mapred-site.xml file: mv mapred-site.xml.template mapred-site.xml
    Close the file.
15. Open /usr/local/hadoop/etc/hadoop/yarn-site.xml:
    **gedit /usr/local/hadoop/etc/hadoop/yarn-site.xml**
    Add these lines inside <configuration> </configuration> tags
    **<property>**
    **<name>yarn.nodemanager.aux-services</name>**
    **<value>mapreduce_shuffle</value>**
    **</property>**
16. Perform following commands to set-up SSH public key authentication
    `ssh-keygen -t rsa -P ""`
    `cat /home/bigdata/.ssh/id_rsa.pub >> /home/bigdata/.ssh/authorized_keys`
    Note: **bigdata** is a user here. It may be different in your case.
17. Format HDFS: **hadoop namenode -format**
18. Start service: **start-all.sh**
19. Execute command to check services are running: **jps**

20. Once your job is done, you can stop services: **`stop-all.sh`**
    Do not execute this command now, first do something in HDFS cluster as described below.

**HDFS command**

1. Assume you are at your home directory or execute: **`cd`**
2. Examine files in HDFS cluster: **`hdfs dfs -ls`**
   You will see nothing in the output as you have not created anything yet.
3. Create a directory in HDFS cluster: **`hdfs dfs -mkdir /user`**
4. Examine files again in HDFS cluster: **`hdfs dfs -ls`**
   You will see the listing of /user directory in the output
5. Create a directory inside /user directory using your user name:
   **`hdfs dfs -mkdir /user/bigdata`**
6. Examine that directory in HDFS cluster: **`hdfs dfs -ls /user/`**
7. Create a text file, say **`demo.txt`**, using gedit and write something in it and save it.
8. Upload the file in the cluster**:**
   **`hdfs dfs -copyFromLocal demo.txt /user/bigdata/`**
9. List the file in the cluster: **`hdfs dfs -ls /user/bigdata/`**
10. See the content of file: **`hdfs dfs -cat /user/bigdata/demo.txt`**